Signal Loom AI

Media In. Machine Intelligence Out. — signalloomai.com

TRANSCRIPTION ID
trn_01hx7kp2q4f8n0r9m3yd

SOURCE FILE	DURATION	SPEAKERS	LANGUAGE
q3-platform-review.mp4	01:47:23	4 identified	English (en)

29	4	3	5	2	97%
SEGMENTS	SPEAKERS	TOPICS	ENTITIES	ACTION ITEMS	ACCURACY

OUTPUT FORMATS DELIVERED

- ✓ TXT Transcript
- ✓ Timestamps
- ✓ Speaker Labels
- ✓ SRT Captions
- ✓ VTT Captions
- ✓ Structured JSON

SPEAKER SUMMARY

- SPEAKER_00 Sarah Chen — confidence 0.97 — 10 turns
- SPEAKER_01 Marcus Rivera — confidence 0.94 — 8 turns
- SPEAKER_02 Dr. A. Patel — confidence 0.91 — 7 turns
- SPEAKER_03 James Whitfield — confidence 0.89 — 4 turns

FULL TRANSCRIPT — TIMESTAMPED & SPEAKER-LABELED

00:29 → 01:04 SPEAKER_00 conf: 0.97 Q3 Platform Review

Alright, let's get started. The goal for today's Q3 review is to walk through what the platform team shipped in July and align on August priorities before the board meeting next week.

01:08 → 01:43 SPEAKER_01 conf: 0.94 PRODUCT: recommendation endpoint

Sounds good. I also want to flag the latency regression we saw on the recommendation endpoint last Tuesday — it affected about twelve percent of users during peak hours.

01:47 → 02:19 **SPEAKER_00** conf: 0.96

Good catch. We've already rolled back the offending commit. The post-mortem is in Notion under eng/post-mortems/2026-Q3. I'll share the link in the follow-up email.

02:22 → 02:51 **SPEAKER_02** conf: 0.91 ML Infrastructure

To add to that — the ML team also shipped the new embedding pipeline ahead of schedule. We're seeing a forty percent reduction in retrieval latency on vector search queries.

03:04 → 03:31 **SPEAKER_03** conf: 0.89 ORG: ML team

That's solid progress. What about the audio transcription pipeline — are we still planning to integrate the new Whisper model by end of quarter?

03:35 → 04:02 **SPEAKER_00** conf: 0.97 DATE: Q4

Yes — the ML infrastructure team has already validated the model. We're targeting a mid-October deployment for that. I'll put James in touch with Dr. Patel to coordinate the testing window.

TOPICS IDENTIFIED

Q3 Platform Review	91%
ML Infrastructure & Embedding Pipeline	88%
Latency Regression — Recommendation Endpoint	85%
Q4 Roadmap Planning	82%

ENTITIES EXTRACTED

PRODUCT	recommendation endpoint	0.89	ORG	ML team	0.87
ORG	platform team	0.85	DATE	last Tuesday	0.82
PERCENT	12% / 40%	0.91	PERSON	James Whitfield, Dr. A. Patel	0.88

ACTION ITEMS DETECTED

⚠ OPEN
SPEAKER_01 — Marcus Rivera
Share post-mortem Notion link for Q3 latency regression
OPEN

✓ CLOSED
SPEAKER_00 — Sarah Chen
Roll back offending commit (recommendation endpoint regression)
CLOSED

SAMPLE — STRUCTURED JSON OUTPUT


```
{
  "id": "trn_01hx7kp2q4f8n0r9m3yd",
  "status": "success",
  "duration_seconds": 6423,
  "speakers": [
    { "id": "spk_00", "label": "Sarah Chen", "confidence": 0.97 },
    { "id": "spk_01", "label": "Marcus Rivera", "confidence": 0.94 }
  ],
  "topics": [
    { "label": "Q3 Platform Review", "confidence": 0.91, "span": [29, 64] }
  ],
  "entities": [
    { "type": "PRODUCT", "text": "recommendation endpoint", "confidence": 0.89 }
  ],
  "action_items": [
    { "speaker": "spk_01", "text": "Share post-mortem Notion link", "status": "open" }
  ]
}
```